
Variational Spectral Graph Convolutional Networks

Louis Tiao*

`louis.tiao@sydney.edu.au`
The University of Sydney

Pantelis Elinas

`pantelis.elinas@data61.csiro.au`
CSIRO's Data61

Harrison Nguyen*

`harrison.nguyen@sydney.edu.au`
The University of Sydney

Edwin V. Bonilla

`edwin.bonilla@data61.csiro.au`
CSIRO's Data61

Abstract

We propose a Bayesian approach to spectral graph convolutional networks (GCNs) where the graph parameters are considered as random variables. We develop an inference algorithm to estimate the posterior over these parameters and use it to incorporate prior information that is not naturally considered by standard GCNs. The key to our approach is to define a smooth posterior parameterization over the adjacency matrix characterizing the graph, which we estimate via stochastic variational inference. Our experiments show that we can outperform standard GCN methods in the task of semi-supervised classification in noisy-graph regimes.

1 Introduction

Graphs represent the elements of a system and their relationships as a set of nodes and edges between them. By exploiting the inter-dependencies of these elements, many applications of machine learning have achieved significant success, for example in the areas of social networks [13], document classification [22, 39] and bioinformatics [11]. Recent notable theoretical and practical insights in graph-based analysis [14, 5, 15, 22] along with breakthroughs in deep learning and big data processing, have reignited research interest in problems such as network analysis and discovery [21], relational learning [37], semi-supervised learning [22, 39], generative models for graphs [7, 25], and, more generally, graph-based inductive inference [13].

In particular, motivated by the incredible success of convolutional neural networks (CNNs) [24] on regular-grid data, researchers have generalized some of their fundamental properties (such as their ability to learn local stationary structures efficiently) to graph-structured data [5, 15, 8]. These approaches mainly focused on exploiting feature dependencies explicitly defined by a graph in an analogous way to how CNNs model long-range correlations through local interactions across pixels in an image. The seminal work by [22] leveraged these ideas to model *dependencies across instances* (instead of features) to be able to incorporate knowledge of the instances' relationships in a semi-supervised learning setting, going beyond the standard i.i.d. assumption.

In this work we focus precisely in the problem of semi-supervised classification based on the method developed in [22], which is now commonly referred to as graph convolutional networks (GCNs). These networks can be seen as a first-order approximation of the spectral graph convolutional networks developed by [8], which itself built upon the pioneering work of [5, 15]. The great popularity of GCNs is mainly due to their practical performance as, at the time it was published, it outperformed related methods by a significant margin. Another practical advantage of using GCNs is their relatively simple propagation rule, which does not require expensive operations such as eigen-decomposition.

*Work done while at CSIRO's Data61.

However, one of the main assumptions underlying GCNs is that the given graph is helpful for the task at hand and that the corresponding edges are highly reliable. This is generally not true in practical applications, as the given graph maybe (i) noisy, (ii) loosely related to the classification problem of interest or (iii) built in an ad hoc basis using e.g. side information. In practice, it is difficult to incorporate this type of uncertainty over the graph using the original GCN framework in a principled way and the performance of the method degrades significantly with increasingly noisy graphs. In this paper we address exactly this problem by adopting a Bayesian approach that places a prior over the graph structure, as given by the adjacency matrix. This, however, poses significant inference challenges as estimating the posterior over the adjacency matrix under a highly non-linear likelihood (as given by the GCN’s output) is intractable.

Thus, to estimate this posterior we resort to approximations as given by stochastic variational inference (VI), which is underpinned by the maximization of the evidence lower bound (ELBO). Nevertheless, a crucial remaining challenge is how to parameterize the posterior over $\mathcal{O}(N^2)$ parameters, where N is the number of nodes in the graph. We propose a simple but effective parameterization inspired by low-rank matrix factorization, where we assume underlying continuous latent variables for each row/column of the posterior parameters over the adjacency matrix. To the best of our knowledge, we are the first to analyze the behavior of graph neural networks under noisy regimes systematically and to develop an efficient variational Bayesian approach that incorporates uncertainty information about the given graph directly. Our experiments show that we can outperform standard GCN methods and other competitive baselines in the task of semi-supervised classification in noisy-graph regimes.

1.1 Related Work

Graph neural networks. Most graph-CNN frameworks can be seen from a more general perspective under the unifying mixture models network (MONET) [29]. Before GCNs, the problem of semi-supervised classification had been traditionally tackled via optimization of a supervised loss along with graph Laplacian regularization [45, 3, 40] or via graph embeddings [32, 38, 42]. GCNs, as proposed by [20] seem to have been the first to use spectral convolutional networks for this problem. However, these approaches are inherently transductive and do not generalize to unseen nodes. Hence, improvements to GCNs include inductive generalizations [13] and ‘attention’ mechanisms [39]. In fact, graph attention networks (GATs) as proposed by [13] can be somewhat more resilient to noisy settings. For more details of graph neural network approaches the reader is referred to the excellent related work in [39] and, more generally, to the surveys in [44, 41].

Uncertain graphs. Previous works have looked at learning settings with uncertain graphs. For example, [4] proposes a new graph-anonymization technique that injects uncertainty in the existence of edges of the graph. [6] proposes a method that models the probability of a node belonging to a particular class as a function of the uncertainty in the edges related to that node. However, such an approach does not either build upon state-of-the-art graph networks or propose a full coherent probabilistic model over the parameters of the network. Following a different methodology, [16] deals with the problem of uncertain graphs via embeddings by constructing a proximity matrix given the uncertain graph and applying matrix factorization to get the embedded representation. The embeddings are then used in supervised learning tasks. Unlike our method, this a two-step procedure where the prediction task is separate from the uncertainty modeling and embedding learning.

Probabilistic approaches. Having a similar motivation to ours, [43] propose a probabilistic model for GCNs where the graph is considered as a realization of mixed membership stochastic block models [2]. However, despite their method being referred to as Bayesian, it parameterizes the *true posterior* over the graph directly and this posterior is not dependent on the observed data (neither features or labels). Thus, it can be seen more like an ensemble GCN. Furthermore, unlike ours, their main focus is on the low-labeled-data regime, although they also present some results in an adversarial setting. Similarly targeting the low-labeled-data regime, [30] proposes a method for semi-supervised classification with Gaussian process (GP) priors, where the parameters of a robust-max likelihood for a node are given by the average of the GP values over its 1-hop neighborhood. Their results indicate that their method can outperform GCNs in active learning settings. Finally, [23] proposes the variational graph autoencoder, a probabilistic framework that also learns latent representations for graphs but, unlike our work, is designed for the task of link prediction.

2 Spectral Graph Convolution Models

Let $\mathbf{X} \in \mathbb{R}^{N \times D}$ be a set of D -dimensional features representing N instances and $\mathbf{Y} = \{\mathbf{y}_n\}$ be their corresponding labels, some of which are observed and others unobserved and $\mathbf{y}_n \in \{0, 1\}^C$ is one-hot-encoded. The goal of semi-supervised classification is to leverage the labeled and unlabeled data in order to predict the unobserved labels. In this paper we are interested in doing so by explicitly exploiting the dependencies among data-points as given by a graph \mathcal{G} .

To do this, we consider the popular GCN models as proposed by [22], which can be seen as first-order approximations to more general (but computationally costly) spectral graph convolutional networks [8]. Let $\mathcal{G} = (\mathcal{V}, \mathcal{E})$ denote an undirected graph with N nodes $v_i \in \mathcal{V}$, edges $(v_i, v_j) \in \mathcal{E}$, binary adjacency matrix $\mathbf{A} \in \{0, 1\}^{N \times N}$, degree matrix \mathbf{D} with $D_{ii} = \sum_j A_{ij}$ and define the normalized graph Laplacian $\mathbf{L} \equiv \mathbf{I}_N - \mathbf{D}^{-\frac{1}{2}} \mathbf{A} \mathbf{D}^{-\frac{1}{2}}$. As proposed in [14] and adopted by [8], spectral graph convolutions can be defined as the multiplication of a signal with a filter in the Fourier domain, which can be well-approximated by a truncated decomposition in terms of Chebyshev polynomials $T_k(\cdot)$, i.e. $g_\theta \star x \approx \sum_{k=0}^K \theta_k T_k(\tilde{\mathbf{L}})$, where θ are Chebyshev coefficients; $\tilde{\mathbf{L}} = \frac{2}{\lambda_0} \mathbf{L} - \mathbf{I}_N$; λ_0 is the largest eigenvalue of \mathbf{L} and \mathbf{I}_N is the N -dimensional identity matrix.

[22] showed that when considering a first order approximation to the above expansion (i.e. $K = 1$), and under a few simplifying assumptions, one can rewrite, for a signal \mathbf{X} , the convolved signal matrix as $\tilde{\mathbf{D}}^{-\frac{1}{2}} \tilde{\mathbf{A}} \tilde{\mathbf{D}}^{-\frac{1}{2}} \mathbf{X} \mathbf{W}$, where \mathbf{W} is a matrix of filter parameters, $\tilde{\mathbf{A}} = \mathbf{A} + \mathbf{I}_N$ is the adjacency matrix of the graph augmented with self-loops and $\tilde{\mathbf{D}}$ is the corresponding degree matrix with $\tilde{D}_{ii} = \sum_{j=1}^N \tilde{A}_{ij}$. This convolved signal matrix constitutes the basic operation in GCNs.

2.1 Composition of graph convolutions for semi-supervised classification

Hence, we can define compositions of these (approximate) spectral graph convolutions $\mathbf{f}^{(l)}(\mathbf{X}, \mathbf{A})$ by the recurrence relation,

$$\begin{aligned} \mathbf{f}^{(0)}(\mathbf{X}, \mathbf{A}) &= \mathbf{X}, \\ \mathbf{f}^{(l+1)}(\mathbf{X}, \mathbf{A}) &= h^{(l+1)} \left(\tilde{\mathbf{D}}^{-\frac{1}{2}} \tilde{\mathbf{A}} \tilde{\mathbf{D}}^{-\frac{1}{2}} \mathbf{f}^{(l)}(\mathbf{X}, \mathbf{A}) \mathbf{W}_l \right), \end{aligned} \quad (1)$$

where $h^{(l)}(\cdot)$ is a nonlinear activation function of the l -th layer, typically the element-wise rectified linear unit (RELU), $\text{relu}(\cdot) = \max(0, \cdot)$.

GCN uses these types of compositions to define a neural network architecture for semi-supervised classification, where the activation for the final layer $h^{(L)}(\cdot)$ is the row-wise softmax function. Each layer is parameterized by \mathbf{W}_l , a $Q^{(l)} \times Q^{(l+1)}$ matrix of weights, where $Q^{(0)} = D$ and $Q^{(L)} = C$. The number of hidden units $Q^{(l)}$ is a positive integer for layers $0 < l < L$.

We will denote the GCN parameters with θ , which consists of the weights for all layers $\theta = \{\mathbf{W}_l\}_{l=1}^L$ and are trained, in the original method, by minimization of the cross-entropy error.

Example: 2-layer GCN. In this paper we will focus on two-layer architectures, hence the output of the GCN is given by:

$$\mathbf{\Pi} = \mathbf{f}^{(L)}(\mathbf{X}, \mathbf{A}) = \text{softmax} \left(\tilde{\mathbf{D}}^{-\frac{1}{2}} \tilde{\mathbf{A}} \tilde{\mathbf{D}}^{-\frac{1}{2}} \text{relu} \left(\tilde{\mathbf{D}}^{-\frac{1}{2}} \tilde{\mathbf{A}} \tilde{\mathbf{D}}^{-\frac{1}{2}} \mathbf{X} \mathbf{W}_0 \right) \mathbf{W}_1 \right), \quad (2)$$

where $\mathbf{\Pi}$ is an $N \times C$ matrix of probabilities for all nodes and classes.

As shown in [22], when the given graph is highly reliable, GCN models trained with the proposed cross-entropy minimization method can yield state-of-the-art classification results. However, as we shall see in § 4, when the given graph is noisy, GCN's performance can degrade significantly and we require encoding our beliefs about the given graph in a principled way as well as having alternative methods to estimate its parameters. In order to address this problem, we propose a Bayesian approach that considers the graph parameters as random variables and uses approximate inference to estimate the posterior over these parameters. Such a posterior is then used in conjunction with the GCN parameters for making predictions over the unlabeled instances.

2.2 Likelihood

Our likelihood model assumes that, conditioned on all the features and the graph adjacency matrix, the observed labels \mathbf{Y}^o are conditionally independent, i.e.,

$$p_{\theta}(\mathbf{Y}^o | \mathbf{X}, \mathbf{A}) = \prod_{\mathbf{y}_n \in \mathbf{Y}^o} p_{\theta}(\mathbf{y}_n | \mathbf{X}, \mathbf{A}) \quad \text{with} \quad p_{\theta}(\mathbf{y}_n | \mathbf{X}, \mathbf{A}) = \text{Cat}(\mathbf{y}_n | \boldsymbol{\pi}_n), \quad (3)$$

where $\text{Cat}(\mathbf{y}_n | \boldsymbol{\pi}_n)$ denotes a Categorical distribution over \mathbf{y}_n with parameters $\boldsymbol{\pi}_n$ being the n -th row of the $N \times C$ probability matrix $\boldsymbol{\Pi}$ obtained from a GCN with L layers, i.e., $\boldsymbol{\Pi} = \mathbf{f}^{(L)}(\mathbf{X}, \mathbf{A})$. As before, θ denotes the GCN’s weight parameters. One of the fundamental differences of our approach with standard GCNs is that we consider a prior over graphs that is constructed using the observed (but potentially noisy or unreliable) adjacency matrix.

2.3 Prior over graphs

We consider random graph priors of the form

$$p(\mathbf{A}) = \prod_{ij} p(A_{ij}), \quad \text{with} \quad p(A_{ij}) = \text{Bern}(A_{ij} | \rho_{ij}^o), \quad (4)$$

where $\text{Bern}(A_{ij} | \rho_{ij}^o)$ is a Bernoulli distribution over A_{ij} with parameter ρ_{ij}^o . Our prior is constructed given an observed (auxiliary) graph $\bar{\mathbf{A}}$ but, for simplicity in the notation, we omit this conditioning here and in all related distributions. This prior can be constructed in various ways so as to encode our beliefs about the graph structure and about how much this structure should be trusted a priori for our semi-supervised classification task. For example, we have found that the construction $\rho_{ij}^o = \bar{\rho}_1 \bar{\mathbf{A}}_{ij} + \bar{\rho}_0 (1 - \bar{\mathbf{A}}_{ij})$, with $0 < \bar{\rho}_1, \bar{\rho}_0 < 1$ being predefined constants, works very well in practice as it gives just enough flexibility to encode our degree of belief on the absence and presence of links separately. Alternatively, we can also consider a hierarchical approach where we place a prior over the parameters of the Bernoulli distribution, for example a Beta distribution. However, we have found experimentally that there are not additional benefits from introducing this hierarchy in the prior. In contrast, as we shall discuss in the next section, and in the experiments, having smooth characterizations of the approximate posterior distribution over the adjacency does yield crucial benefits in inference.

3 Graph structure inference via smooth parameterizations

Our goal is to carry out joint inference over the GCN parameters and the graph structure as given by the adjacency matrix. Since the main additional component of our approach is to consider a prior over the adjacency matrix, we focus on the estimation of the posterior over this matrix given the observed data, i.e. $p(\mathbf{A} | \mathbf{X}, \mathbf{Y}^o) \propto p_{\theta}(\mathbf{Y}^o | \mathbf{X}, \mathbf{A}) p(\mathbf{A})$ where the likelihood and prior terms are given by Equations eqs. 3 and 4, respectively.

Computation of this posterior is analytically intractable due to the highly-nonlinear nature of the likelihood so we resort to approximate posterior inference methods. Given the high-dimensional nature of the posterior over \mathbf{A} , we focus on compact representations of the posterior via VI [18]. Generally, VI methods entail the definition of an approximate posterior (variational) distribution and the estimation of its parameters via the maximization of the so-called evidence lower bound (ELBO), a procedure that is known to be equivalent to minimizing the Kullback-Leibler (KL) divergence between the approximation posterior and the true posterior.

3.1 Variational distribution: free vs smooth parameterizations

Similarly to the prior definition, our approximate posterior is of the form

$$q_{\phi}(\mathbf{A}) = \prod_{ij} q(A_{ij}), \quad \text{with} \quad q_{\phi}(A_{ij}) = \text{Bern}(A_{ij} | \rho_{ij}), \quad \rho_{ij} > 0, \quad (5)$$

where, henceforth, we use ϕ to denote all the parameters of the variational posterior. In the case where ρ_{ij} are free parameters then $\phi = \{\rho_{ij}\}$. We refer to this approach as the ‘free’ parameterization.

Unfortunately, as illustrated in the supplement (appendix D.1), such a parameterization makes optimization of the ELBO wrt ϕ extremely difficult and one is forced to use alternative representations of the posterior. Intuitively, conditional independence in the posterior is a strong assumption and small changes in ρ_{ij} will compete with each other to explain the data. Consequently, any continuous optimization algorithm will find it very challenging to find a good direction in this non-smooth combinatorial space. Therefore, we propose the following smooth parameterization:

$$\rho_{ij} = \sigma(\mathbf{z}_i^T \mathbf{z}_j + b_i + b_j + s), \quad \mathbf{z}_i \in \mathbb{R}^d, \quad \{b_i, s \in \mathbb{R}\}, \quad i, j = 1, \dots, N, \quad (6)$$

where $\sigma(x) \equiv (1 + \exp(-x))^{-1}$ is the logistic sigmoid function and $d \leq D$ is the dimensionality of the latent representation \mathbf{z} . As we see, the same representation is shared across the columns and rows of \mathbf{A} 's Bernoulli parameters, which addresses the combinatorial nature of the optimization landscape of the free parameterization. We note that this parameterization is referred to in the matrix-factorization and link-prediction literature as low-rank [28, 27] or dot-product [23]. In this case the variational parameters are $\phi = \{\{\mathbf{z}_i, b_i\}, s\}$.

3.2 Variational distribution: discrete vs relaxed

We have defined above a variational distribution which naturally models the discrete nature of the adjacency matrix \mathbf{A} . Our goal is to estimate the parameters ϕ of the posterior $q_\phi(\mathbf{A})$ via maximization of the ELBO. For this purpose we can use the so-called score function method [33], which provides an unbiased estimator of the gradient of an expectation of a function using Monte Carlo (MC) samples. However, it is now widely accepted that, because of its generality, the score function estimator can suffer from high variance [34].

Therefore, as an alternative to the score function estimator, we can use the so-called re-parameterization trick [20, 35], which generally exhibits lower variance. Unfortunately, the re-parameterization trick is not applicable to discrete distributions and we need to resort to continuous relaxations. In this work we use Concrete distributions as proposed by [17, 26]. In particular, we denote our binary Concrete posterior distribution with location parameters $\lambda_{ij} > 0$ and temperature $\tau > 0$ as $q_\phi(A_{ij}) = \text{BinConcrete}(A_{ij} | \lambda_{ij}, \tau)$. Analogously, as discussed in [26], in order to maintain a lower bound during variational inference we also relax our prior so that $p(A_{ij}) = \text{BinConcrete}(A_{ij} | \lambda_{ij}^o, \tau_o)$. In this case the variational parameters are the parameters of the Concrete distribution which can be, as in the discrete case, free parameters $\phi = \{\lambda_{ij}\}$ or have a smooth parameterization analogous to that in eq. 6, i.e. $\lambda_{ij} = \sigma(\mathbf{z}_i^T \mathbf{z}_j + b_i + b_j + s)$ and, consequently, $\phi = \{\{\mathbf{z}_i, b_i\}, s\}$.

3.3 Evidence Lower Bound

It is easy to show that we can write the ELBO as

$$\mathcal{L}_{\text{ELBO}}(\phi) = \mathbb{E}_{q_\phi(\mathbf{A})} \log p_\theta(\mathbf{Y}^o | \mathbf{X}, \mathbf{A}) - \text{KL}[q_\phi(\mathbf{A}) \parallel p(\mathbf{A})], \quad (7)$$

where $\mathbb{E}_{q_\phi(\mathbf{A})} \log p_\theta(\mathbf{Y}^o | \mathbf{X}, \mathbf{A})$ is the expected log likelihood (ELL), i.e. the expectation of the conditional likelihood over the approximate posterior, and $\text{KL}[q_\phi(\mathbf{A}) \parallel p(\mathbf{A})]$ is the KL divergence between the approximate posterior and the prior. When using discrete (Bernoulli) prior and posterior distributions the KL term can be determined analytically and the ELL is estimated via MC samples. When using the binary Concrete relaxations for the prior and the posterior, one requires sampling from the approximate posterior and evaluation of $\log q_\phi(A_{ij})$. These can be done straightforwardly and details can be found in the supplement. As described above we can use the score-function estimator or the re-parameterization trick within gradient-based optimization of the ELBO when we have discrete distributions or their continuous relaxations, respectively. For completeness, here we give the full expression of the ELBO when using the Concrete relaxations under a more numerically stable parameterization (see supplement for details):

$$\mathcal{L}_{\text{ELBO}}(\phi) = \mathbb{E}_{g_{\phi, \tau}(\mathbf{B})} \left[\log p_\theta(\mathbf{Y}^o | \mathbf{X}, \sigma(\mathbf{B})) - \log \frac{g_{\phi, \tau}(\mathbf{B})}{f_{\tau_o}(\mathbf{B})} \right], \quad \text{where} \quad (8)$$

$$g_{\phi, \tau}(B_{ij}) = \text{Logistic}(B_{ij} | \frac{\log \lambda_{ij}}{\tau}, \frac{1}{\tau}), \quad f_{\tau_o}(B_{ij}) = \text{Logistic}(B_{ij} | \frac{\log \lambda_{ij}^o}{\tau_o}, \frac{1}{\tau_o}), \quad (9)$$

$\sigma(\mathbf{B})$ computes the entrywise logistic sigmoid function over \mathbf{B} ; $\text{Logistic}(B | \mu, s)$ denotes a Logistic distribution with location μ and scale s and the distributions $g_{\phi, \tau}(\mathbf{B})$ and $f_{\tau_o}(\mathbf{B})$ factorize over the

entries of \mathbf{B} . The expectation $E_{g_{\phi, \tau}(\mathbf{B})}$ is estimated using S samples from the re-parameterized posterior, which can be obtained using eq. 11 below. Estimation of the variational parameters ϕ is done via gradient-based optimization of the ELBO with the gradients obtained by automatic differentiation.

3.4 Predictions

The posterior predictive distribution of the latent labels, given our factorized assumption of the conditional likelihood in eq. 3, can be obtained as

$$p(\mathbf{Y}^u | \mathbf{Y}^o, \mathbf{X}) = \sum_{\mathbf{A}} p_{\theta}(\mathbf{Y}^u | \mathbf{X}, \mathbf{A}) p(\mathbf{A} | \mathbf{Y}^o, \mathbf{X}) \approx \frac{1}{S} \sum_{s=1}^S p_{\theta}(\mathbf{Y}^u | \mathbf{X}, \mathbf{A}^{(s)}), \quad (10)$$

where $\mathbf{A}^{(s)}$ is a sample from the posterior $q_{\phi}(\mathbf{A})$, S is the total number of samples and $p_{\theta}(\mathbf{Y}^u | \mathbf{X}, \mathbf{A})$ is the GCN-likelihood given in eq. 3. These samples can be obtained as:

$$U \sim \text{Uniform}(0, 1), \quad A_{ij}^{(s)} = \sigma \left(\frac{\log \lambda_{ij} + \log U - \log(1 - U)}{\tau} \right), \quad (11)$$

where, as before, $\{\lambda_{ij}\}$ are the estimated parameters of the posterior and τ is the posterior temperature.

3.5 Implementation and computational complexity

We implement our approach using TensorFlow [1] for efficient GPU-based computation and also use some components of TensorFlow Probability [9]. We will release our code and datasets if this paper gets accepted. The time complexity of our algorithm can be derived from considering the two main components of the ELBO in eq. 7, namely the KL divergence term and the ELL term. We focus here on one-hidden layer GCN (apart from the output layer) with dimensionality $Q \equiv Q^{(1)}$ along with a smooth (dot product) parameterization of the posterior. We recall that N is the number of nodes, D is the dimensionality of the input features \mathbf{X} , d is the dimensionality of the latent variables \mathbf{Z} used to parameterize the posterior, C is the number of classes and S is the number of samples from the variational posterior used to estimate the required expectations.

KL divergence term: For both the discrete and relaxed parameterization we need to compute the dot-product parameterization over each A_{ij} which is $\mathcal{O}(d)$ and, given these parameters, $\mathcal{O}(N^2)$ individual KL divergences, whose computation can be trivially parallelized but gradient information must be aggregated for each \mathbf{z}_i . While for a discrete posterior these individual KL terms can be computed exactly (as shown in the supplement), for the continuous relaxation we need to resort to MC estimation over S samples. Aggregation over samples can also be parallelized straightforwardly.

ELL term: Computing the ELL using a 2-layer GCN as in eq. 2 requires $\mathcal{O}(NDQ + S(NQC + N^2Q + N^2C))$ for the continuous case. However, in the discrete case it only requires doing a forward pass over the standard GCN architecture S times, hence being linear in the number of edges, i.e. $\mathcal{O}(S|\mathcal{E}|DQC)$, where $|\mathcal{E}|$ is the expected number of edges sampled from the posterior, assuming sparse-dense matrix multiplication is exploited.

4 Experiments

In this section we describe the experiments carried out to evaluate our Bayesian approach and compare with several benchmark models. Throughout this section and in the supplement, we will refer to our method as Bayesian graph convolutional network (BGCN) and its relaxed and discrete variants as BGCN-R and BGCN-D, respectively. We will first focus on real graphs that have been corrupted by a noise process that adds *fake links*. This is motivated by the fact that in many practical applications graphs are created in an ad hoc basis based on side information or node features [15]. Then we present an experiment where in fact the graphs are created by this very same process. Finally, we analyze the behavior of our algorithms on a real noisy dataset on Twitter data. All the results presented here are based on the relaxed version of our algorithm, which we found to outperform the discrete version consistently.

Table 1: Datasets used in the experiments. Label rate refers to the proportion of labeled nodes used for training.

Dataset	Type	Nodes	Edges	Classes	Features	Label rate
CORA	Citation network	2,708	5,429	7	1,433	0.052
CITeseer	Citation network	3,327	4,732	6	3,703	0.036
TWITTER	Social network	4,971	14,591	2	205	0.01

Datasets: We use two citation-network datasets (CORA and CITESEER) and another real dataset based on Twitter data (TWITTER). In the citation networks the nodes represent documents and their links refer to citations between documents. We construct an undirected graph based on these citation links so as to obtain a binary adjacency matrix. Each document is characterized by a set of features, which are either bag-of-words (BOW) or term frequency-inverse document frequency (TF-IDF) for CORA and CITESEER, respectively. The details of these datasets are shown in [table 1](#). The class for each document is given by their subject and our goal is to predict this for a subset of documents in the network.

TWITTER is based on the dataset published in [36]. Each node represents a Twitter user. An edge between two users exists if a user has re-tweeted the other user. Each node has an associated feature vector such that each feature is derived either from the user’s activity or the user’s tweet texts. Description of these features is given in the supplement ([appendix D.5](#)). The class for each user is given by a binary status property that indicates whether the user is hateful or not; our goal is to predict a user’s hatefulness status. We note that the Twitter graph is derived from users’ activities and it is not equal to the Twitter social network based on follower relationships. Because of this, we believe this dataset is a good example of a real noisy graph.

Experimental set-up: For the citation networks we used the fixed splits as in [42] where 20 labeled examples per class are used for training, 1,000 examples are used for testing and the rest is used as unlabeled data. We compare against standard GCN [22], sample-and-aggregate (GRAPHSAGE) [13] and GAT [39]. Throughout our experiments we use the test accuracy as given by the proportion of correctly classified test examples and the test mean log likelihood (MLL). In both cases, the higher the better. Full details of the experimental set-up including parameter settings and a description of the experiments on TWITTER are given in the supplement ([appendix C](#)).

Noisy graphs (fake links): Here we analyze the performance of the proposed method and the baselines’ when the original graph is corrupted via the insertion of fake links. These are added to the graph as a proportion of the number of existing links, which we refer to as corruption factor. for example, a corruption factor of 0.5 in CORA indicates that around 2,700 fake edges have been added. Also notice that even a corruption factor of 6 represents a very small proportion of all the potential links in the graph. In CORA this amounts to around 0.07%. Accuracy results (MLL values are reported in the supplement) for CITESEER are shown in [fig. 1](#) (left) where we can see that our method, for all the smoothing parameters considered, outperforms the baselines significantly after considering a corruption factor greater than 2. The advantages of our method are less pronounced in CORA ([fig. 2](#), right) but we still see significant improvements for higher noise levels. We note that in the noiseless case, i.e. when using the original uncorrupted graph, our method achieves similar performance to GCN with an appropriate smoothing factor $\bar{\rho}_1 > 0.9$. Similarly, when the graph is not used (corresponding to very small smoothing factors, $\bar{\rho}_1 = \bar{\rho}_0 = 10^{-5}$), our method performs similarly to GCN, which is equivalent to using a standard multi-layer Perceptron neural network. The results of these settings are shown in the supplement ([appendix D.3](#)).

Feature-based graphs: Here we consider feature-based graphs, i.e. graphs that have been constructed using node features, in the same spirit as in [15]. This is common practice across practitioners as in many realistic scenarios a graph is not given yet we want to exploit the dependencies across instances with state-of-the-art methods such as GCN, GRAPHSAGE or GAT. These graphs are built using an algorithm based on K -nearest neighbors, see the supplement for details. In this experiment we examine the performance of our method based on two parameters: the number of hidden units in the GCN likelihood ($Q \equiv Q^{(1)}$) and the number of neighbors used to construct the graph (K). The results are shown in [fig. 2](#). As in the noisy case, we see that our method consistently improves performance over the baselines on CITESEER (left) and CORA (right), with the exception of two




Figure 1: Results for noisy graphs on CITESEER (left) and CORA (right): accuracy as a function of the ratio of additional fake links relative to the number of true links. BGCN-R- $\bar{\rho}_1$ stands for the relaxed version of our method with smoothing factor $\bar{\rho}_1$.

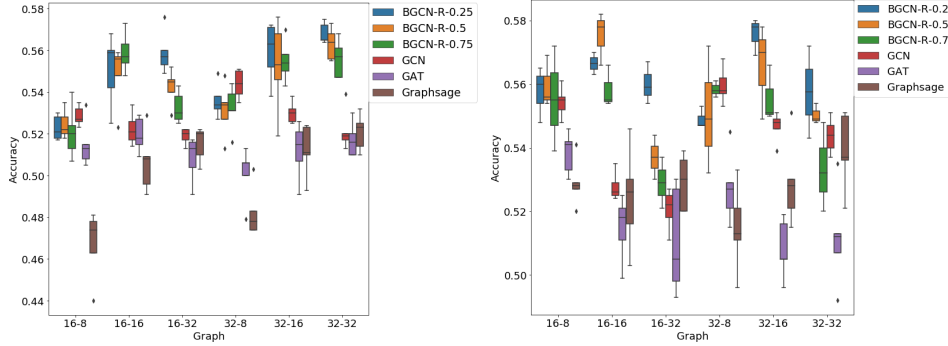


Figure 2: Results for feature-based graphs on CITESEER (left) and CORA (right): Accuracy as a function of Q - K settings, where Q is the number of hidden units and K the number of neighbors used to construct the graph. BGCN-R- $\bar{\rho}_1$ stands for the relaxed version of our method with smoothing factor $\bar{\rho}_1$.

$Q - K$ combinations. As before, the benefits of our approach are more pronounced for CITESEER than for CORA. We attribute these differences to the types of features used, BOW vs TF-IDF.

Twitter data: Results on TWITTER are analyzed in the supplement (appendix D.5) where we see that, on average, our method can perform similarly or better than standard GCNs. However, in this particular experiment all methods seem to exhibit high variance due to the large differences in performance across different data splits.

5 Conclusion & discussion

We have presented a Bayesian approach to spectral graph convolutional networks by building upon state-of-the-art neural network architectures for graph-based semi-supervised classification. We have shown that the performance of standard GCN architectures can degrade significantly when considering noisy graphs. Our method deals with this problem in a natural way by considering prior distributions over the adjacency matrix along with a GCN likelihood in a joint probabilistic model. Our inference algorithm for posterior estimation relies on a smooth parameterization of the parameters of the posterior and has been shown to work very well in real citation graph datasets. An immediate direction for future work is to address the scalability of our method, as the relaxed posterior version cannot exploit sparse computations. A related avenue is to improve upon the discrete

posterior version of our algorithm as it does exploit sparse operations but uses high-variance gradient estimates.

A Binary discrete distributions

The KL divergence between two Bernoulli distributions $q(a | \rho)$ and $p(a | \rho^o)$ can be computed as

$$\text{KL}[q(a | \rho) \parallel p(a | \rho^o)] = \rho[\log \rho - \log \rho^o] + (1 - \rho)[\log(1 - \rho) - \log(1 - \rho^o)]. \quad (12)$$

B Binary concrete distributions

In this section we give details of the re-parameterization used for the implementation of our algorithm when both the prior and the approximate posterior are relaxed via the binary Concrete distribution [26, 17].

B.1 Summary of Bernoulli relaxation transformations

$$A_{ij} \sim \text{BinConcrete}(\lambda_{ij}, \tau) \quad \Leftrightarrow \quad A_{ij} = \sigma(B_{ij}), \quad B_{ij} \sim \text{Logistic}\left(\frac{\log \lambda_{ij}}{\tau}, \frac{1}{\tau}\right) \quad (13)$$

$$B_{ij} \sim \text{Logistic}\left(\frac{\log \lambda_{ij}}{\tau}, \frac{1}{\tau}\right) \quad \Leftrightarrow \quad B_{ij} = \frac{\log \lambda_{ij} + L}{\tau}, \quad L \sim \text{Logistic}(0, 1) \quad (14)$$

$$L \sim \text{Logistic}(0, 1) \quad \Leftrightarrow \quad L = \sigma^{-1}(U) := \log U - \log(1 - U) \quad U \sim \text{Uniform}(0, 1) \quad (15)$$

In summary, we have

$$A_{ij} \sim \text{BinConcrete}(\lambda_{ij}, \tau) \quad \Leftrightarrow \quad A_{ij} = \sigma\left(\frac{\log \lambda_{ij} + \sigma^{-1}(U)}{\tau}\right) \quad U \sim \text{Uniform}(0, 1). \quad (16)$$

B.2 Re-parameterized ELBO

With the results above, it is easy to see that we can write the ELBO as:

$$\mathcal{L}_{\text{ELBO}}(\phi) = \mathbb{E}_{q_{\phi, \tau}(\mathbf{A})} \left[\log p_{\theta}(\mathbf{Y} | \mathbf{X}, \mathbf{A}) - \log \frac{q_{\phi, \tau}(\mathbf{A})}{p_{\tau_o}(\mathbf{A})} \right] \quad (17)$$

$$= \mathbb{E}_{g_{\phi, \tau}(\mathbf{B})} \left[\log p_{\theta}(\mathbf{Y} | \mathbf{X}, \sigma(\mathbf{B})) - \log \frac{q_{\phi, \tau}(\sigma(\mathbf{B}))}{p_{\tau_o}(\sigma(\mathbf{B}))} \right] \quad (18)$$

$$= \mathbb{E}_{g_{\phi, \tau}(\mathbf{B})} \left[\log p_{\theta}(\mathbf{Y} | \mathbf{X}, \sigma(\mathbf{B})) - \log \frac{g_{\phi, \tau}(\mathbf{B})}{f_{\tau_o}(\mathbf{B})} \right]. \quad (19)$$

where

$$g_{\phi, \tau}(B_{ij}) = \text{Logistic}\left(B_{ij} \mid \frac{\log \lambda_{ij}}{\tau}, \frac{1}{\tau}\right), \quad f_{\tau_o}(B_{ij}) = \text{Logistic}\left(B_{ij} \mid \frac{\log \lambda_{ij}^o}{\tau_o}, \frac{1}{\tau_o}\right). \quad (20)$$

C Full details of experimental set up

For standard GCN and our methods we use a two-layer GCN as given in eq. 2. We train standard GCN exactly as done by [22] so as to minimize the cross-entropy loss on the same datasets, i.e. using dropout, L2 regularization, Glorot weight initialization [12] and row-normalization of input-feature vectors. Hyperparameters were set to those found to perform best by [22], i.e. dropout rate of 0.5, L2 regularization of 5×10^{-4} and 16 hidden units. The ADAM optimizer [19] was used with an initial learning rate of 0.01.

For our method, we carry out point estimation of the GCN-likelihood parameters and the approximate posterior parameters so as to maximize the ELBO in eq. 7. Hyperparameters for GCN-likelihood estimation were the same as above. To construct the prior over the adjacency matrix we followed




Figure 3: The negative evidence lower bound (NELBO) as a function of the number of epochs for the freely parameterized posterior (Free) and the smooth parameterization (Lowrank).

the procedure explained in § 2.3 with various smoothing parameters $\bar{\rho}_1$ and $\bar{\rho}_0 = 10^{-5}$. Posterior initialization was done to be as close as possible to the prior by setting the latent variables \mathbf{Z} via low-rank factorization with the number of latent dimensions set to $d = 1,000$. The biases $\{b_i\}$ and shift s in eq. 6 were initialized to 0 and -20 , respectively. Prior and posterior temperatures τ_o and τ were set to 0.1. All the methods were executed up to a maximum of 5,000 epochs (unless otherwise stated explicitly) with $S = 3$ samples used for estimating the required expectations.

We used our own implementation for GAT and GRAPHSAGE as provided by the open source graph machine learning library STELLARGRAPH². GAT used an architecture identical to the one described in [39]. The first layer consists of $K = 8$ attention heads computing $F = 8$ features, followed by an exponential linear unit (ELU) nonlinearity. The second layer is used for classification that computes C features (where C is the number of classes), followed by a softmax activation. $L2$ regularization with $\lambda = 0.0005$ and 0.6 dropout was used. The implementation of GRAPHSAGE used mean aggregator functions and sampled the neighborhood at a depth of $K = 2$ with neighborhood sample size of $S_1 = S_2 = 32$. The model was trained with 0.5 dropout and $L2$ regularization with $\lambda = 0.0005$.

D Additional results

D.1 Smooth parameterization

Figure 3 compares a typical run of our algorithms using the free and smooth parameterizations of the posterior. We see that the optimizer struggles to find directions of improvement for the free parameterization, whereas for the smooth parameterization the negative evidence lower bound (NELBO) decreases steadily during optimization. This behavior was observed throughout several learning rate settings and we attribute it to the discrete combinatorial optimization nature of the NELBO when using the free parameterization.

D.2 Noiseless graphs

Table 2 shows the results for the noisy graphs where we see that our approach performs similarly to standard GCNs.

D.3 Noisy graphs

Test mean log likelihood (MLL) results are given in fig. 4.

²<https://www.stellargraph.io/>.

Table 2: Results for noiseless graphs: ‘no graph’ corresponds to ignoring the input graph and ‘full graph’ corresponds to using the full original graph without corruption. Graph convolutional network (GCN); Discrete Bayesian graph convolutional network (BGCN-D) and Relaxed Bayesian graph convolutional network (BGCN-R). Accuracy (ACC) and mean log likelihood (MLL).

Method	CORA- no graph				CORA- full graph			
	ACC	(std dev)	MLL	(std dev)	ACC	(std dev)	MLL	(std dev)
GCN	0.5592	0.0119	-1.3058	0.0309	0.8178	0.0041	-0.6287	0.0107
BGCN-D	0.5494	0.0084	-1.3290	0.0222	0.8049	0.0024	-0.6760	0.0092
BGCN-R	0.5584	0.0102	-1.2985	0.0257	0.8193	0.0045	-0.6306	0.0076

Method	CITESEER- no graph				CITESEER- full graph			
	ACC	(std dev)	MLL	(std dev)	ACC	(std dev)	MLL	(std dev)
GCN	0.5537	0.0151	-1.2474	0.0166	0.7041	0.0068	-0.9342	0.0075
BGCN-D	0.5484	0.0082	-1.2654	0.0105	0.6837	0.0053	-1.0351	0.0222
BGCN-R	0.5491	0.0071	-1.2631	0.0122	0.7056	0.0041	-0.9418	0.0072




Figure 4: Results for noisy graphs on CITESEER (left) and CORA (right): Mean log likelihood (MLL) as a function of the ratio of additional fake links relative to the number of true links. BGCN-R- $\bar{\rho}_1$ stands for the relaxed version of our method with smoothing factor $\bar{\rho}_1$.

D.4 Feature-based graphs

Following the methods of [15], a graph was estimated using CORA and CITESEER based on the node features. Given a training set with features and labels we train a fully connected network with 1-hidden layer, $Q \in \{16, 32\}$ neurons using standard RELU activation and 0.5 dropout between each layer. We then extract the first layer features $\tilde{\mathbf{Z}}_1 \in \mathbb{R}^{M \times Q}$ where M includes training, validation and test sets and consider their distance, $d(i, j) = \|\tilde{\mathbf{Z}}_{1,i} - \tilde{\mathbf{Z}}_{1,j}\|^2$. The adjacency matrix, \mathbf{A} , was constructed by taking the $K \in \{8, 16, 32\}$ closest neighbors. Test mean log likelihood (MLL) results are given in fig. 5.

D.5 Twitter

The TWITTER dataset we use here is a subset of that published in [36]³. We modified the original dataset in two ways. First, the graph in [36] has 100,000 nodes of which only 4,971 have manually been annotated as hateful or not. For our experiments we consider the subgraph of the annotated nodes only. The original graph has approximately 2,200,000 edges whereas the annotated users subgraph has 14,591 edges. Second, we use a subset of the node features in [36]. The latter includes manually engineered graph features such as several node centrality measures. Furthermore, [36]

³This dataset can be downloaded from Kaggle at <https://www.kaggle.com/manoelribeiro/hateful-users-on-twitter>




Figure 5: Results for feature-based graphs on CITESEER (left) and CORA (right): Mean log likelihood (MLL) as a function of Q - K settings, where Q is the number of hidden units and K the number of neighbors used to construct the graph. BGCN-R- $\bar{\rho}_1$ stands for the relaxed version of our method with smoothing factor $\bar{\rho}_1$.

includes features derived from the lexical analysis of the users’ tweets based on two different text representation learning methods, Glove [31] and Empath [10]. We have removed all the graph-based manually engineered features. Lastly, we only consider the Empath lexical features and ignore the Glove ones since we expect them to encode similar information about the content of a user’s tweets. The nodes in the TWITTER dataset for our empirical study have feature vectors with 205 dimensions.⁴

We split the data into training, validation, and test sets as follows. We use 25 labeled examples from each class for training such that the size of the training set is 50 or only 1% of the total number of nodes in the graph. We then split the remaining nodes in to validation and test sets such that the validation set contains 30% of the data and the test set the remaining 70%. All data are sampled uniformly at random. We generated 20 such splits of the data and used them consistently in the evaluation.

We compare 4 Bayesian GCN models, the relaxed version, using different values for the edge smoothing parameter and prior, posterior temperatures against a GCN and an MLP model. All BGCN, GCN, and MLP models use a single hidden layer with 16 hidden units. The experimental setup is as described here in [appendix C](#) and in the main paper [§ 4](#) unless stated otherwise.

The results are shown in Figure 6. For our method the notation BGCN- $\bar{\rho}_1$ -Pr-Po indicates that $\bar{\rho}_1$ is the smoothing factor for the existing edges in the graph; Pr is the temperature used for the prior distribution; and Po is the temperature used for the posterior distribution. We trained all models for a maximum 5000 epochs and used the accuracy on the validation set for early stopping. We report the accuracy and mean log likelihood on the test set for the latter model. Each model was trained and evaluated on each of the 20 splits of the data. The results in Figure 6 indicate that all models exhibit high predictive variance. Generally, BGCN models with a high degree of smoothing and temperature parameters in the range suggested in [26], show a trend towards better performance with regards to the mean or median values for each set of experiments. However, we emphasize that the high predictive variance of all models indicates that we cannot say that any model is statistically significantly better than the rest.

References

- [1] Martín Abadi, Ashish Agarwal, Paul Barham, et al. TensorFlow: Large-scale machine learning on heterogeneous systems, 2015.
- [2] Edoardo M Airolidi, David M Blei, Stephen E Fienberg, and Eric P Xing. Mixed membership stochastic blockmodels. *Journal of Machine Learning Research*, 9(Sep):1981–2014, 2008.

⁴We will be releasing this dataset together with the source code implementing our variational spectral graph convolutional network algorithm.




Figure 6: Results for TWITTER: Box plots for machine learning model vs accuracy (left) and mean log likelihood (right) on the TWITTER test set. BGCN- $\bar{\rho}_1$ -Pr-Po indicates that $\bar{\rho}_1$ is the smoothing factor for the existing edges in the graph; Pr is the temperature used for the prior distribution; and Po is the temperature used for the posterior distribution. For each box plot, the vertical lines indicate median values and the triangles indicate the mean values of the data. Outliers are shown using diamonds.

- [3] Mikhail Belkin, Partha Niyogi, and Vikas Sindhwani. Manifold regularization: A geometric framework for learning from labeled and unlabeled examples. *Journal of Machine Learning Research*, 7(Nov):2399–2434, 2006.
- [4] Paolo Boldi, Francesco Bonchi, Aristides Gionis, and Tamir Tassa. Injecting uncertainty in graphs for identity obfuscation. *Proceedings of the VLDB Endowment*, 5(11):1376–1387, 2012.
- [5] Joan Bruna, Wojciech Zaremba, Arthur Szlam, and Yann Lecun. Spectral networks and locally connected networks on graphs. In *International Conference on Learning Representations*, 2014.
- [6] Michele Dallachiesa, Charu Aggarwal, and Themis Palpanas. Node classification in uncertain graphs. In *International Conference on Scientific and Statistical Database Management*. ACM, 2014.
- [7] Nicola De Cao and Thomas Kipf. Molgan: An implicit generative model for small molecular graphs. *arXiv preprint arXiv:1805.11973*, 2018.
- [8] Michaël Defferrard, Xavier Bresson, and Pierre Vandergheynst. Convolutional Neural Networks on Graphs with Fast Localized Spectral Filtering. In *Advances in Neural Information Processing Systems*, pages 3844–3852, 2016.
- [9] Joshua V Dillon, Ian Langmore, Dustin Tran, Eugene Brevdo, Srinivas Vasudevan, Dave Moore, Brian Patton, Alex Alemi, Matt Hoffman, and Rif A Saurous. TensorFlow Distributions. *arXiv preprint arXiv:1711.10604*, 2017.
- [10] Ethan Fast, Binbin Chen, and Michael S Bernstein. Empath: Understanding topic signals in large-scale text. In *Proceedings of the 2016 CHI Conference on Human Factors in Computing Systems*, pages 4647–4657. ACM, 2016.
- [11] Alex Fout, Jonathon Byrd, Basir Shariat, and Asa Ben-Hur. Protein interface prediction using graph convolutional networks. In *Advances in Neural Information Processing Systems*, pages 6530–6539, 2017.
- [12] Xavier Glorot and Yoshua Bengio. Understanding the difficulty of training deep feedforward neural networks. In *Artificial Intelligence and Statistics*, pages 249–256, 2010.
- [13] Will Hamilton, Zhitao Ying, and Jure Leskovec. Inductive representation learning on large graphs. In *Advances in Neural Information Processing Systems*, pages 1024–1034, 2017.
- [14] David K Hammond, Pierre Vandergheynst, and Rémi Gribonval. Wavelets on Graphs via Spectral Graph Theory. *Applied and Computational Harmonic Analysis*, 30(2):129–150, 2011.
- [15] Mikael Henaff, Joan Bruna, and Yann LeCun. Deep convolutional networks on graph-structured data. *arXiv preprint arXiv:1506.05163*, 2015.
- [16] Jiafeng Hu, Reynold Cheng, Zhipeng Huang, Yixang Fang, and Siqiang Luo. On embedding uncertain graphs. In *Proceedings of the 2017 ACM on Conference on Information and Knowledge Management*, pages 157–166. ACM, 2017.

- [17] Eric Jang, Shixiang Gu, and Ben Poole. Categorical Reparameterization with Gumbel-Softmax. *arXiv preprint arXiv:1611.01144*, 2016.
- [18] Michael I Jordan, Zoubin Ghahramani, Tommi S Jaakkola, and Lawrence K Saul. An introduction to variational methods for graphical models. *Machine learning*, 37(2):183–233, 1999.
- [19] Diederik P Kingma and Jimmy Ba. Adam: A method for stochastic optimization. *arXiv preprint arXiv:1412.6980*, 2014.
- [20] Diederik P Kingma and Max Welling. Auto-encoding variational Bayes. In *International Conference on Learning Representations*, 2014.
- [21] Thomas Kipf, Ethan Fetaya, Kuan-Chieh Wang, Max Welling, and Richard Zemel. Neural relational inference for interacting systems. *arXiv preprint arXiv:1802.04687*, 2018.
- [22] Thomas N Kipf and Max Welling. Semi-Supervised Classification with Graph Convolutional Networks. *arXiv preprint arXiv:1609.02907*, 2016.
- [23] Thomas N Kipf and Max Welling. Variational graph auto-encoders. *arXiv preprint arXiv:1611.07308*, 2016.
- [24] Yann LeCun, Léon Bottou, Yoshua Bengio, Patrick Haffner, et al. Gradient-based learning applied to document recognition. *Proceedings of the IEEE*, 86(11):2278–2324, 1998.
- [25] James Lloyd, Peter Orbanz, Zoubin Ghahramani, and Daniel M Roy. Random function priors for exchangeable arrays with applications to graphs and relational data. In *Advances in Neural Information Processing Systems*, pages 998–1006, 2012.
- [26] Chris J Maddison, Andriy Mnih, and Yee Whye Teh. The Concrete Distribution: A Continuous Relaxation of Discrete Random Variables. *arXiv preprint arXiv:1611.00712*, 2016.
- [27] Aditya Krishna Menon and Charles Elkan. Link prediction via matrix factorization. In *Joint european conference on machine learning and knowledge discovery in databases*, pages 437–452. Springer, 2011.
- [28] Andriy Mnih and Ruslan R Salakhutdinov. Probabilistic matrix factorization. In *Advances in neural information processing systems*, pages 1257–1264, 2008.
- [29] Federico Monti, Davide Boscaini, Jonathan Masci, Emanuele Rodola, Jan Svoboda, and Michael M Bronstein. Geometric deep learning on graphs and manifolds using mixture model cnns. In *IEEE Conference on Computer Vision and Pattern Recognition*, pages 5115–5124, 2017.
- [30] Yin Cheng Ng, Nicolò Colombo, and Ricardo Silva. Bayesian semi-supervised learning with graph gaussian processes. In *Advances in Neural Information Processing Systems*, pages 1683–1694, 2018.
- [31] Jeffrey Pennington, Richard Socher, and Christopher Manning. Glove: Global vectors for word representation. In *Proceedings of the 2014 conference on empirical methods in natural language processing (EMNLP)*, pages 1532–1543, 2014.
- [32] Bryan Perozzi, Rami Al-Rfou, and Steven Skiena. Deepwalk: Online learning of social representations. In *International Conference on Knowledge Discovery and Data Mining*, pages 701–710, New York, NY, USA, 2014. ACM.
- [33] Rajesh Ranganath, Sean Gerrish, and David M. Blei. Black box variational inference. In *Artificial Intelligence and Statistics*, 2014.
- [34] Rajesh Ranganath, Dustin Tran, and David Blei. Hierarchical variational models. In *International Conference on Machine Learning*, pages 324–333, 2016.
- [35] Danilo J Rezende, Shakir Mohamed, and Daan Wierstra. Stochastic backpropagation and approximate inference in deep generative models. In *International Conference on Machine Learning*, 2014.
- [36] Manoel Horta Ribeiro, Pedro H Calais, Yuri A Santos, Virgílio AF Almeida, and Wagner Meira Jr. "Like sheep among wolves": Characterizing hateful users on Twitter. *arXiv preprint arXiv:1801.00317*, 2017.
- [37] Michael Schlichtkrull, Thomas N Kipf, Peter Bloem, Rianne Van Den Berg, Ivan Titov, and Max Welling. Modeling relational data with graph convolutional networks. In *European Semantic Web Conference*, pages 593–607. Springer, 2018.

- [38] Jian Tang, Meng Qu, Mingzhe Wang, Ming Zhang, Jun Yan, and Qiaozhu Mei. Line: Large-scale information network embedding. In *International Conference on World Wide Web*, pages 1067–1077, 2015.
- [39] Petar Veličković, Guillem Cucurull, Arantxa Casanova, Adriana Romero, Pietro Liò, and Yoshua Bengio. Graph Attention Networks. In *International Conference on Learning Representations*, 2018.
- [40] Jason Weston, Frédéric Ratle, Hossein Mobahi, and Ronan Collobert. Deep learning via semi-supervised embedding. In *Neural Networks: Tricks of the Trade*, pages 639–655. Springer, 2012.
- [41] Zonghan Wu, Shirui Pan, Fengwen Chen, Guodong Long, Chengqi Zhang, and Philip S Yu. A comprehensive survey on graph neural networks. *arXiv preprint arXiv:1901.00596*, 2019.
- [42] Zhilin Yang, William W Cohen, and Ruslan Salakhutdinov. Revisiting semi-supervised learning with graph embeddings. In *International Conference on Machine Learning*, 2016.
- [43] Yingxue Zhang, Soumyasundar Pal, Mark Coates, and Deniz Üstebay. Bayesian graph convolutional neural networks for semi-supervised classification. In *AAAI Conference on Artificial Intelligence*, 2019.
- [44] Jie Zhou, Ganqu Cui, Zhengyan Zhang, Cheng Yang, Zhiyuan Liu, and Maosong Sun. Graph neural networks: A review of methods and applications. *arXiv preprint arXiv:1812.08434*, 2018.
- [45] Xiaojin Zhu, Zoubin Ghahramani, and John D Lafferty. Semi-supervised learning using gaussian fields and harmonic functions. In *International Conference on Machine Learning*, pages 912–919, 2003.